

CryptoGuard: Lightweight Hybrid Detection and Response to Host-based Cryptojackers in Linux Cloud Environment

Gyeonghoon Park, Jaehan Kim, Jinu Choi, and Jinwoo Kim




UCIRVINE
UNIVERSITY of CALIFORNIA • IRVINE



광운대학교
KwangWoon University

The Growing Threat of Host-based Cryptojackers

- Target Linux-based public cloud environments
 - In 2018, hackers enlisted Tesla's public cloud to mine cryptocurrency¹
 - In 2023, leaked AWS credentials were used to create EC2 instances for cryptomining²
- Cryptojacking incidents have increased by 659%³




1 <https://www.wired.com/story/cryptojacking-tesla-amazon-cloud/>

2 <https://thehackernews.com/2023/10/elektra-leak-cryptojacking-attacks.html>

3 SonicWall, Cyber Threat Report 2024

The Lifecycle of a Host-based Cryptojacker⁴

Benign Application




- ***Host-based Cryptojacking***
 - Exploits PCs and IoTs
 - Pervasive in modern cloud environments

- ***In-browser Cryptojacking:***
 - Exploits client web browsers connected to malicious websites
 - Slowed after CoinHive shutdown in 2019

Evasion Techniques of Host-based Cryptojackers

- CPU throttling
 - Makes difficult to determine cryptojackers using *rule-based detection* that relies on CPU usage



Evasion Techniques of Host-based Cryptojackers

- **PID obfuscation**
 - Continuously obfuscates PIDs with short and unexpected time intervals

PID	USER	CPU%	MEM%	TIME+	Command	Δ
3879	user1	0.6	0.1	0:00.07	/cryptojacker.bin	

- **Restoration via entry points**
 - Can restore cryptojackers even after a system reboot or process termination by compromising entry points (e.g., cronjob⁵, rc.local⁶)

```
GNU_nano 6.2                               /tmp/crontab.w8a6t5/crontab
* * * * * /home/user1/.cache/mesa_shader_cache/18/b4fqq09
```

Existing Solutions

- **Non-ML detection**
 - Lachtar et al. [IEEE CAL '20], D. Tanana and G. Tanana [IEEE CAL '20]
- **ML detection**
 - Gomes et al. [NCA '20], Caprolu et al. [Comput. Commun. '21],
Tekiner et al. [NDSS 2022]
 - Mani et al. [ACSOS '20]
- **Prevention solution**
 - Franco et al. [IEEE ICC '23] → Suricata IDS alerts

None of them focuses on both detection
and persistent prevention system

Challenges

- How to minimize ***overhead*** when collecting fine-grained features for detection?
 - Network traffic, CPU, HPC, syscall, etc...
- How to counter ***evasion techniques*** in detection?
 - CPU throttling, PID obfuscation, entry point, etc...
- How to achieve ***scalability*** in the cloud environment?




Cryptoguard

- Sketch/sliding window-based syscall monitoring via eBPF⁷
- Precise detection for stealthy cryptojackers via deep learning
- Integrated detection and remediation approaches




CryptoGuard Overview




System Call Monitoring

- Q. Is it practical to monitor system calls in cloud environment?
→ Monitoring all processes and containers imposes significant overhead
- Key idea: Perform monitoring in two phases
 - Phase 1: Host-level monitoring
 - Phase 2: Process-level monitoring




1st Phase: Host-level Monitoring/Detection

- Key idea: **Count-Min Sketch (CMS)**⁸
 - A probabilistic data structure for frequency estimation
 - Record hashed values of system calls into a fixed-size 2D array
- Perform binary classification: suspicious/benign



2nd Phase: Process-level Monitoring/Detection


- Key idea: *Sliding Window*
 - Trace syscalls within a time window of ΔT from the detection point
 - Recent traces are sufficient due to the distinctiveness of patterns
- Perform multi-class classification: xmrig/sysrv/benign



Cryptojacker Remediation

```
openat(AT_FDCWD, "/home/c  
0_RDONLY|0_CLOEXEC) = 484  
→ Check if the cache directory is a  
openat(AT_FDCWD, "/home/c  
nvdmld", 0_WRONLY|0_CREAT|0_TRUNC|0_CLOEXEC, 0777) = 484  
→ Creates a new binary using a randomly generated alphanumeric string
```

Job Schedule_1 Before CryptoGuard:
* * * * * /home/cryptoguard/.cache/mesa_shader_cache/73/nvdmld
Deleted Job Schedule_1 After CryptoGuard:
* * * * * /home/cryptoguard/.cache/mesa_shader_cache/73/nvdmld



Dataset & Implementation

- Collected a dataset by executing both malware samples and benign processes
 - 123 real-world Linux cryptojacking malware samples⁹ were used
 - Benign processes include Apache HTTP Server, Apache Tomcat, Nginx, Redis, MySQL

Family	# of Samples	Class
Sysrv	100	Sysrv
XMRig	15	XMRig
TeamTNT	5	
WatchDog	3	
Total	123	



- 5,000 lines of C code (libbpf and bpftrace)

Performance of Detecting Suspicious Hosts (1st Phase)

- Binary classification to distinguish suspicious hosts from benign ones
- LSTM and CNN classifiers achieved average F1-scores of 96.42% and 95.82%

Sketch Size	Class	LSTM			CNN		
		Precision	Recall	F1-score	Precision	Recall	F1-score
272 × 3	Benign	94.62%	98.88%	96.70%	97.36%	94.39%	95.85%
	Malware	98.74%	94.00%	96.31%	94.18%	97.25%	95.69%
55 × 3	Benign	95.82%	96.21%	96.51%	95.43%	95.28%	95.36%
	Malware	95.84%	96.51%	96.71%	94.75%	94.91%	94.83%
55 × 5	Benign	95.67%	97.63%	96.64%	94.77%	98.88%	96.78%
	Malware	97.32%	95.11%	96.20%	98.74%	94.17%	96.40%



Performance of Detecting Cryptojacker Processes (2nd Phase)

- For $\Delta T = 30$, the CNN classifier achieved F1-scores of 95.62%, 98.54%, and 98.87%

Sliding Window Size	Class	LSTM			CNN		
		Precision	Recall	F1-score	Precision	Recall	F1-score
$\Delta T = 30$	XMRig	75.00%	92.14%	82.69%	97.76%	93.57%	95.62%
	Sysrv	91.50%	90.91%	91.21%	98.70%	98.38%	98.54%
	Benign	98.10%	93.66%	95.83%	98.36%	99.40%	98.87%
$\Delta T = 60$	XMRig	69.11%	87.63%	77.27%	91.40%	87.63%	89.47%
	Sysrv	90.85%	89.68%	90.26%	99.34%	96.77%	98.04%
	Benign	95.76%	89.68%	92.62%	96.89%	98.94%	97.91%
$\Delta T = 90$	XMRig	69.11%	87.63%	77.27%	91.40%	87.63%	89.47%
	Sysrv	90.85%	89.68%	90.26%	99.34%	96.77%	98.04%
	Benign	95.76%	89.68%	92.62%	96.89%	98.94%	97.91%

Evasion Resilience / Overhead

- F1-scores of the three approaches under different CPU usage rates of 12.5%, 25%, 50%, and 75%
 - CryptoGuard offers resilience to CPU throttling evasion attacks
- CPU usage variation during monitoring
 - CryptoGuard imposes minimal overhead (0.29% → 0.35%)



Conclusion

- Host-based Cryptojacker
 - Stealthy behavior
 - Employs obfuscation techniques
 - The vast number of hosts and processes in cloud environments makes detection difficult
- **CryptoGuard**: a lightweight solution for detecting and remediating host-based cryptojacking in cloud environment
 - Sketch-based and sliding window-based monitoring
 - Counters the persistence mechanism by Cryptojacker
 - <https://github.com/PGHOON/CryptoGuard>

Thank you for listening

(hoonp2@uci.edu)

ACKNOWLEDGMENTS: This work was partly supported by the National Research Foundation of Korea (NRF) grant funded by the Korea government (MSIT) (No. RS-2024-00457937, Design and implementation of security layers for secure WebAssembly-based serverless environments, 50%) and the Institute of Information & Communications Technology Planning & Evaluation (IITP) grant funded by the Korea government (MSIT) (No. 2021-0-00118, RS-2021-II210118, Development of decentralized consensus composition technology for large-scale nodes, 50%).